

```
In [1]: import numpy as np
import pandas as pd
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
import seaborn as sns
from sklearn import svm
from sklearn.linear_model import LogisticRegression

In [2]: loan_dataset = pd.read_csv("D:\Machine Learning\Loan Status\Loan_status.csv")

In [3]: loan_dataset

Out[3]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	360
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360
...
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71.0	360
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40.0	180
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253.0	360
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187.0	360
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133.0	360

614 rows x 13 columns

```
In [4]: loan_dataset.shape
Out[4]: (614, 13)

In [5]: loan_dataset.describe()

Out[5]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	614.000000	614.000000	592.000000	600.000000	564.000000
mean	5403.459283	1621.245798	146.412162	342.000000	0.842199
std	6109.041673	2926.248369	85.587325	65.12041	0.364878
min	150.000000	0.000000	9.000000	12.00000	0.000000
25%	2877.500000	0.000000	100.000000	360.00000	1.000000
50%	3812.500000	1188.500000	128.000000	360.00000	1.000000
75%	5795.000000	2297.250000	168.000000	360.00000	1.000000
max	81000.000000	41667.000000	700.000000	480.00000	1.000000

```
In [6]: loan_dataset.dtypes

Out[6]:
```

Loan_ID	object
Gender	object
Married	object
Dependents	object
Education	object
Self_Employed	object
ApplicantIncome	int64
CoapplicantIncome	float64
LoanAmount	float64
Loan_Amount_Term	float64
Credit_History	float64
Property_Area	object
Loan_Status	object
dtype:	object

```
In [7]: loan_dataset.isnull().sum()

Out[7]:
```

Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0
dtype:	int64

```
In [8]: #Missing Value

In [9]: loan_dataset['Gender'].fillna(loan_dataset['Gender'].mode()[0], inplace=True)
loan_dataset['Married'].fillna(loan_dataset['Married'].mode()[0], inplace=True)
loan_dataset['Self_Employed'].fillna(loan_dataset['Self_Employed'].mode()[0], inplace=True)
loan_dataset['Credit_History'].fillna(loan_dataset['Credit_History'].mode()[0], inplace=True)
loan_dataset['Dependents'].fillna(loan_dataset['Dependents'].mode()[0], inplace=True)

In [10]: loan_dataset.isnull().sum()

Out[10]:
```

Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	0
Property_Area	0
Loan_Status	0
dtype:	int64

```
In [11]: loan_dataset['Loan_Amount_Term'].value_counts()

Out[11]:
```

360.0	512
180.0	44
480.0	15
300.0	13
84.0	4
240.0	4
120.0	3
36.0	2
60.0	2
12.0	1

Name: Loan_Amount_Term, dtype: int64

```
In [12]: loan_dataset['Loan_Amount_Term'].fillna(loan_dataset['Loan_Amount_Term'].mode()[0], inplace=True)

In [13]: loan_dataset['Dependents'].value_counts()

Out[13]:
```

0	360
1	102
2	101
3+	51

Name: Dependents, dtype: int64

```
In [14]: loan_dataset['LoanAmount'].fillna(loan_dataset['LoanAmount'].median(), inplace=True)

In [15]: loan_dataset.isnull().sum()

Out[15]:
```

Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	0
Property_Area	0
Loan_Status	0
dtype:	int64

```
In [16]: loan_dataset = loan_dataset.replace(to_replace='3+', value=4)

In [17]: loan_dataset['Dependents'].value_counts()

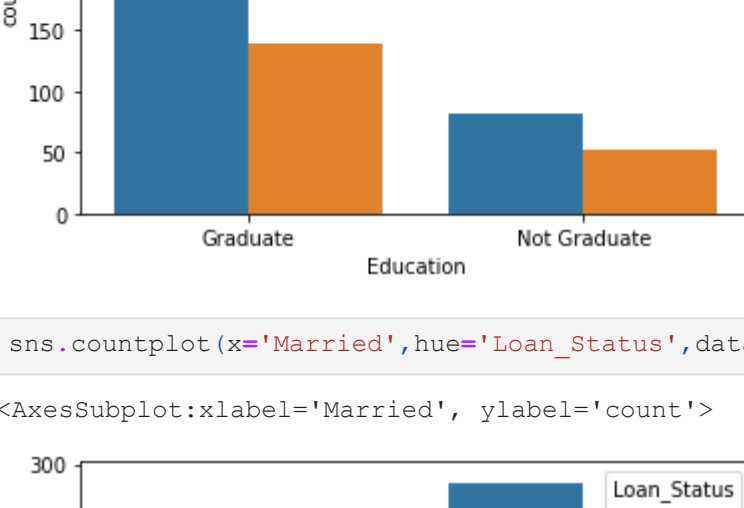
Out[17]:
```

0	360
1	102
2	101
4	51

Name: Dependents, dtype: int64

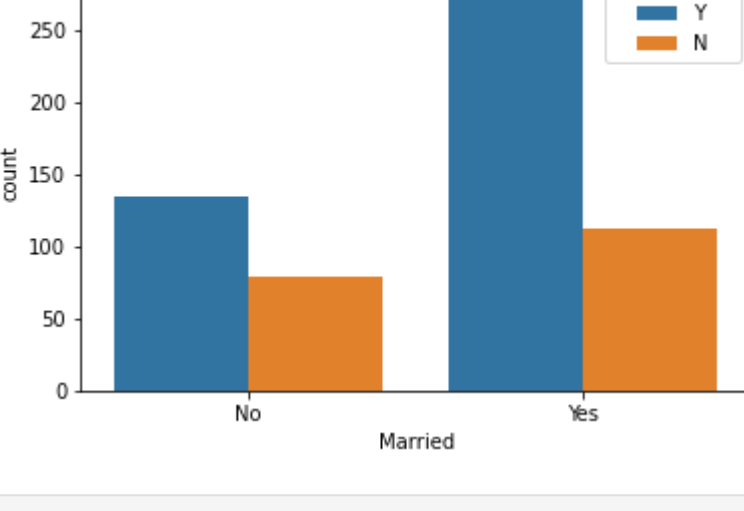
```
In [18]: sns.countplot(x='Education', hue='Loan_Status', data=loan_dataset)

Out[18]: <AxesSubplot:xlabel='Education', ylabel='count'>
```



```
In [19]: sns.countplot(x='Married', hue='Loan_Status', data=loan_dataset)

Out[19]: <AxesSubplot:xlabel='Married', ylabel='count'>
```



```
In [20]: loan_dataset.replace({'Married':{'No':0,'Yes':1}, 'Gender':{'Male':1,'Female':0}, 'Self_Employed':{'No':0,'Yes':1}, 'Property_Area':{'Rural':0,'Semiurban':1,'Urban':2}, 'Education':{'Graduate':1,'Not Graduate':0}})

In [21]: loan_dataset.head()

Out[21]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term
0	LP001002	1	0	0	1	0	5849	0.0	128.0	360.0
1	LP001003	1	1	1	1	0	4583	1508.0	128.0	360.0
2	LP001005	1	1	0	1	1	3000	0.0	66.0	360.0
3	LP001006	1	1	0	0	0	2583	2358.0	120.0	360.0
4	LP001008	1	0	0	1	0	6000	0.0	141.0	360.0

```
In [22]: X = loan_dataset.drop(columns=['Loan_ID','Loan_Status'],axis=1)
Y = loan_dataset['Loan_Status']

In [23]: X

Out[23]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
0	1	0	0	1	0	5849	0.0	128.0	360.0	0
1	1	1	1	1	0	4583	1508.0	128.0	360.0	1
2	1	1	0	1	1	3000	0.0	66.0	360.0	1
3	1	1	0	0	0	2583	2358.0	120.0	360.0	1
4	1	0	0	1	0	6000	0.0	141.0	360.0	1
...
609	0	0	0	1	0	2900	0.0	71.0	360.0	0
610	1	1	4	1	0	4106	0.0	40.0	180.0	1
611	1	1	1	1	0	8072	240.0	253.0	360.0	1
612	1	1	2	1	0	7583	0.0	187.0	360.0	1
613	0	0	0	1	1	4583	0.0	133.0	360.0	0

614 rows x 11 columns

```
In [24]: Y

Out[24]:
```

0	Y
1	N
2	Y
3	Y
4	Y
...	...
609	Y
610	Y
611	Y
612	Y
613	N

Name: Loan_Status, Length: 614, dtype: object

```
In [25]: X_train, X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.1,stratify=Y,random_state=2)

In [26]: X.shape

Out[26]: (614, 11)

In [27]: X_train.shape

Out[27]: (552, 11)

In [28]: X_test.shape

Out[28]: (62, 11)

In [29]: #classifier=svm.SVC(kernel='linear')

In [30]: model = LogisticRegression()

In [31]: model.fit(X_train,Y_train)

C:\Users\GAURAV RATHOD\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:762: ConvergenceWarning: 1
bfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
LogisticRegression()

In [32]: # accuracy score on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction,Y_train)

In [33]: print('Accuracy on training data : ', training_data_accuracy)

Accuracy on training data : 0.8079710144927537

In [34]: X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction,Y_test)

In [35]: print('Accuracy on training data : ', test_data_accuracy)

Accuracy on training data : 0.7741935483870968

In [37]:
```

```
from sklearn import svm
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV

model_params = {
    'svm': {
        'model': svm.SVC(gamma='auto'),
        'params' : {
            'C': [1,10,20],
            'kernel': ['rbf','linear']
        }
    },
    'random_forest': {
        'model': RandomForestClassifier(),
        'params': {
            'n_estimators': [1,5,10,15,20]
        }
    },
    'logistic_regression': {
        'model': LogisticRegression(solver='liblinear',multi_class='auto'),
        'params': {
            'C': [1,15,30]
        }
    },
    'DecisionTreeClassifier': {
        'model':DecisionTreeClassifier(random_state=0),
        'params':{
            'max_depth':[3,4,5,6,7,8,9]
        }
    },
    'KNN':{
        'model':KNeighborsClassifier(),
        'params':{
            'n_neighbors':[3,5,7,9,11,13,15]
        }
    }
}
scores = []

for model_name, mp in model_params.items():
    clf = GridSearchCV(mp['model'], mp['params'], cv=5, return_train_score=False)
    clf.fit(X_test, Y_test)
    scores.append({
        'model': model_name,
        'best_score': clf.best_score_,
        'best_params': clf.best_params_
    })

df = pd.DataFrame(scores,columns=['model','best_score','best_params'])
df
```

```
Out[37]:
```

	model	best_score	best_params
0	svm	0.743590	{'C': 20, 'kernel': 'linear'}
1	random_forest	0.710256	{'n_estimators': 20}
2	logistic_regression	0.743590	{'C': 1}
3	DecisionTreeClassifier	0.666667	{'max_depth': 4}
4	KNN	0.693590	{'n_neighbors': 13}

```
In [ ]:
```